Chirag Agarwal

Website: chirag126.github.io
Email: chiragagarwall12@gmail.com
Phone: (865)-406-2887

Research Interests

TrustworthyML for Foundation Models. While there has been remarkable progress in developing large-scale complex neural network models in recent years, our understanding of how, what, and why they learn what they learn remains shallow. I approach these questions through the lens of interpretability, explainability, robustness, fairness, and privacy. Examining these trustworthy properties will advance our understanding of foundation models such as large language and vision-language models.

Understanding and Improving Model Training Dynamics. Current work on TrustworthyML focuses on developing new explanation methods, but there is little research on using them to understand model training dynamics or designing architectures that incorporate explanations. My research aims to develop techniques to understand the training dynamics of DNN and human interpretable foundation models and use explanations to identify more difficult examples during training for safety-critical applications.

Academic & Professional Experience

Harvard University
Postdoctoral Research Fellow   2020 – Present
Host: Prof. Hima Lakkaraju and Prof. Marinka Zitnik

Adobe
Research Scientist   2022 – 2023

Auburn University
Research Assistant   Summer 2019

Robert Bosch LLC
Computer Vision/Augmented Reality Intern   Summer 2018

Tempus labs Inc.
Imaging Science Intern   Spring 2018

Kitware Inc.
Research and Development Intern   Summer 2017

Geisinger Health Systems
Research Intern   Summer 2016

Education

University of Illinois at Chicago
Ph.D. in Electrical and Computer Engineering   2020
- Thesis: *Robustness and Explainability of Deep Neural Networks*
- Committee: Prof. Dan Schonfeld, Prof. Bharati Prasad, Prof. Mojtaba Soltanalian, Prof. Piotr Gmytrasiewicz, Prof. Anh Nguyen

University of Illinois at Chicago
M.S. in Electrical and Computer Engineering   2018

Selected Honors & Achievements

Top Reviewer for NeurIPS	2023
Spotlight presentation, NeurIPS Ro-FoMo Workshop in Foundation Models	2023
Spotlight paper, ICML	2021
AINet Fellow by DAAD	2021
Spotlight presentation, ICML workshop on Human Interpretability in Machine Learning	2020
Spotlight paper, IEEE Conference on Image Processing (ICIP)	2019
Finalist for the Deans Scholarship Award at UIC	2019

Selected Grants & Awards

Adobe Data Science Research Award (US $50,000) – co-PI	2023
Harvard Data Science Initiative Microsoft Azure Credits (US $22,224) – co-PI	2023
Amazon Research Award (US $70,000) (Under review) – co-PI	2023
NSF Small: Privacy Issues in Language Models Prof. Hima Lakkaraju, Prof. Seth Neel, and Chirag Agarwal (US $600,000) (Under preparation)	2023
AI for Social Good Google Workshop (US $10,000) – co-PI	2021
2 × Research Proposal accepted by Google Cloud Platform (US $1,000) – Sole PI	2020

Research Articles

† denotes the author I co-mentored with the PI; * indicates an equal contribution.

Articles in Peer-Reviewed Journals

49. **C. Agarwal**, O. Queen†, H. Lakkaraju, M. Zitnik: Evaluating Explainability for Graph Neural Networks, *Nature Scientific Data, 2023.*
    128+ GitHub stars

48. H. Honarvar, **C. Agarwal**, S. Somani, A. Vaid, J. Lampert, T. Wanyan, V. Y. Reddy, G. N. Nadkarni, R. Miotto1, M. Zitnik, F. Wang, B. S. Glicksberg: Enhancing convolutional neural network predictions of electrocardiograms with left ventricular dysfunction using a novel sub-waveform representation, *Cardiovascular Digital Health Journal, 2022.*

47. **C. Agarwal**, S. Gupta†, M. Y. Najjar, T. E. Weaver, X. J. Zhou, D. Schonfeld, B. Prasad: Deep Learning Analyses of Brain MRI to Identify Sleepiness in Treated Obstructive Sleep Apnea: A Pilot Study, *Journal of Sleep and Vigilance (JSV), 2022.*

46. B. Prasad*, **C. Agarwal***, E. Schonfeld, D. Schonfeld, B. Mokhlesi: Deep learning applied to polysomnography to predict blood pressure in obstructive sleep apnea and obesity hypoventilation: A proof-of-concept study, *Journal of Clinical Sleep Medicine (JCSM), 2020.*

45. **C. Agarwal**, J. Klobusicky, D. Schonfeld: Convergence of backpropagation with momentum for network architectures with skip connections, *Journal of Computational Mathematics (JCM), 2019.*

44. E. Cha, Y. Veturi, **C. Agarwal**, M. Arbabshirani, S. Pendergrass: Using Adipose Measures from Electronic Health Record Imaging Based Data for Discovery, *Journal of Obesity, 2018.*

Articles in Peer-Reviewed Conference Proceedings

43. S. H. Tanneru†, **C. Agarwal**, H. Lakkaraju: Uncertainty In Explanations Of Large Language Models, *International Conference on Artificial Intelligence and Statistics (AISTATS), 2024.*
    Spotlight Presentation at the NeurIPS R0-FoMo Workshop, 2023

42. M. Llordes, D. Ganguly, S. Bhatia, **C. Agarwal**: Explain like I am BM25: Interpreting a Dense Model's Ranked-List with a Sparse Approximation, *ACM SIGIR Conference on Research and Development in Information Retrieval (SIGIR), 2023.*

41. A. Seth, M. Hemani, **C. Agarwal**: DeAR: Debiasing Vision-Language Models with Additive Residuals, *Conference on Computer Vision and Pattern Recognition (CVPR), 2023.*

40. S. Deshmukh†, A. Dasgupta, B. Krishnamurthy, N. Jiang, **C. Agarwal**, J. Subramanian, G. Theocharous: Trajectory-based Explainability Framework for Offline RL, *International Conference on Learning Representations (ICLR), 2023.*

39. J. Cheng†, G. Dasoulas, H. He, **C. Agarwal**, M. Zitnik: GNNDelete: A General Unlearning Strategy for Graph Neural Networks, *International Conference on Learning Representations (ICLR), 2023.*

38. V. Giunchiglia, C. V. Shukla, G, Gonzalez, **C. Agarwal**: Towards Training GNNs using Explanation Directed Message Passing, *Proceedings of the First Learning on Graphs Conference (LoG), 2022.*

37. **C. Agarwal**, E. Saxena†, S. Krishna†, M. Pawelczyk†, N. Johnson†, I. Puri†, M. Zitnik, H. Lakkaraju: OpenXAI: Towards a Transparent Evaluation of Model Explanations, *Conference on Neural Information Processing Systems (NeurIPS), 2022.*
193+ GitHub stars

36. **C. Agarwal**, D. D'Souza†, S. Hooker: Estimating Example Difficulty using Variance of Gradients, *Conference on Computer Vision and Pattern Recognition (CVPR), 2022.*
57+ GitHub stars

35. **C. Agarwal**, M. Zitnik, H. Lakkaraju: Probing GNN Explainers: A Rigorous Theoretical and Empirical Analysis of GNN Explanation Methods, *International Conference on Artificial Intelligence and Statistics (AISTATS), 2022.*

34. M. Pawelczyk†, **C. Agarwal**, S. Joshi, S. Upadhyay, H. Lakkaraju: Exploring Counterfactual Explanations Through the Lens of Adversarial Examples: A Theoretical and Empirical Analysis, *International Conference on Artificial Intelligence and Statistics (AISTATS), 2022.*

33. **C. Agarwal**, H. Lakkaraju, M. Zitnik: Towards a Unified Framework for Fair and Stable Graph Representation Learning, *Conference on Uncertainty in Artificial Intelligence (UAI), 2021.*

32. S. Agarwal, S. Jabbari, **C. Agarwal**, S. Upadhyay, Z. S. Wu, H. Lakkaraju: Towards the Unification and Robustness of Perturbation and Gradient Based Explanations, *International Conference on Machine Learning (ICML), 2021.*
Spotlight Presentation

31. **C. Agarwal***, S. Khobahi*, D. Schonfeld, M. Soltanalian: CoroNet: A Deep Network Architecture for Semi-Supervised Task-Based Identification of COVID-19 from Chest X-ray Images, *SPIE Medical Imaging, 2021.*

30. **C. Agarwal**, A. Nguyen: Explaining image classifiers by removing input features using generative models, *Asian Conference on Computer Vision (ACCV), 2020.*

29. N. Bansal*, **C. Agarwal***, A. Nguyen*: SAM: The Sensitivity of Interpretability Methods to Hyperparameters, *Conference on Computer Vision and Pattern Recognition (CVPR), 2020.*
Oral presentation (Top 5%)

28. **C. Agarwal**, S. Khobahi, A. Bose, M. Soltanalian, D. Schonfeld: Deep-URL: A Model-Aware Approach To Blind Deconvolution Based On Deep Unfolded Richardson-Lucy Network, *IEEE Conference on Image Processing (ICIP), 2020.*

27. **C. Agarwal**, A. Nguyen, D. Schonfeld: Improving Robustness to Adversarial Examples by Encouraging Discriminative Features, *IEEE Conference on Image Processing (ICIP), 2019.*
Spotlight presentation (Top 10%)

26. M. Aloraini, M. Sharifzadeh, **C. Agarwal**, D. Schonfeld: Statistical Sequential Analysis for Object-based Video Forgery Detection, *Electronic Imaging, 2019.*

25. N. Khobragade*, **C. Agarwal***: Multi-class segmentation of neuronal electron microscopy images using deep learning, *SPIE Medical Imaging, 2018.*

24. **C. Agarwal**, M. Sharifzadeh, D. Schonfeld: CrossEncoders: A complex neural network compression framework, *IS&T International Symposium on Electronic Imaging, 2018.*

23. M. Sharifzadeh, **C. Agarwal**, M. Aloraini, D. Schonfeld: Convolutional neural network steganalysis's application to steganography, *IEEE Visual Communications and Image Processing (VCIP), 2017.*

22. **C. Agarwal**, A.H. Dallal, M.R. Arbabshirani, A. Patel, G. Moore: Unsupervised quantification of abdominal fat from CT images using Greedy Snakes, *SPIE Medical Imaging, 2017.*

21. A.H. Dallal, **C. Agarwal**, M.R. Arbabshirani, A. Patel, G. Moore: Automatic estimation of heart boundaries and cardiothoracic ratio from chest X-ray images, *SPIE Medical Imaging, 2017.*

20. M.R. Arbabshirani, A.H. Dallal, **C. Agarwal**, A. Patel, G. Moore: Accurate segmentation of lung fields on chest radiographs using deep convolutional networks, *SPIE Medical Imaging, 2017.*

19. **C. Agarwal**, A. Bose, S. Maiti, N. Islam, S.K. Sarkar: Enhanced data hiding method using DWT based on Saliency model, *IEEE International Conference on Signal Processing, Computing and Control (ISPCC), 2013.*

18. S. Maiti, **C. Agarwal**, A. Bose, S.K. Sarkar: Robust data hiding technique in wavelet domain using saliency map, *International Journal of Advances in Engineering and Technology, 2013.*

17. N. Islam S. Maiti, A. Bose, **C. Agarwal**, S. K. Sarkar: An Improved Method of Pre-Filter Based Image Watermarking in DWT Domain, *International Journal of Computer Science and Technology, 2013.*

Preprints and Workshop Articles

16. N. Kroeger†, D. Ley†, S. Krishna†, **C. Agarwal**, H. Lakkaraju: Are Large Language Models Post Hoc Explainers?, *Preliminary version presented at the NeurIPS XAIA Workshop, 2023.*

15. A. Kumar, **C. Agarwal**, S. Srinivas, S. Feizi, H. Lakkaraju: Certifying LLM Safety against Adversarial Prompting, *arXiv, 2023.*

14. A. Java, S. Jandial, **C. Agarwal**: Towards Fair Knowledge Distillation using Student Feedback, *Preliminary version presented at the Efficient Systems for Foundation Models, ICML 2023.*

13. S.V. Deshmukh, Srivatsan R, S. Vijay, J. Subramanian, **C. Agarwal**: Counterfactual Explanation Policies in RL, *Preliminary version presented at "Could it have been different?" Counterfactuals in Minds and Machines Workshop, ICML 2023.*

12. T. R. Menta†, S. Jandial†, A. Patil, Vimal KB, S. Bachu, B. Krishnamurthy, V. N. Balasubramanian, **C. Agarwal**, M. Sarkar: Towards Estimating Transferability using Hard Subsets, *arXiv, 2023.*

11. **C. Agarwal**: Intriguing Properties of Visual-Language Model Explanations, *Preliminary version presented at RTML Workshop, ICLR 2023.*

10. S. Krishna†, **C. Agarwal**, H. Lakkaraju: On the Impact of Adversarially Robust Models on Algorithmic Recourse, *Preliminary version presented at Trustworthy and Socially Responsible ML Workshop, NeurIPS 2022.*

9. **C. Agarwal**, N. Johnson†, M. Pawelczyk†, S. Krishna†, E. Saxena†, M. Zitnik, H. Lakkaraju: Rethinking Stability for Attribution-based Explanations, *Preliminary version presented at PAIR^{2}Struct Workshop, ICLR, 2022.* **Oral Presentation**

8. D. D'Souza†, Z. Nussbaum†, **C. Agarwal**, S. Hooker: A Tale Of Two Long Tails, *Preliminary version presented at Uncertainty & Robustness in Deep Learning Workshop, ICML, 2021.*

7. H. Honarvar, **C. Agarwal**, S. Somani, A. Vaid, J. Lampert, T. Wanyan, V. Y. Reddy, G. N. Nadkarni, R. Miotto1, M. Zitnik, F. Wang, B. S. Glicksberg: A novel representation of electrocardiogram waveforms for enhancing deep learning predictions, *Preliminary version presented at Interpretable Machine Learning in Healthcare Workshop, ICML, 2021.*

6. **C. Agarwal***, P. Chen*, A. Nguyen: Intriguing generalization and simplicity of adversarially trained neural networks, *Preliminary version presented at Human Interpretability in Machine Learning Workshop, ICML, 2020.* **Spotlight Presentation**

5. **C. Agarwal**, B. Dong, D. Schonfeld, A. Hoogs: An explainable adversarial robustness metric for deep learning neural networks, 2018.

4. M. Sharifzadeh, **C. Agarwal**, M. Salarian, D. Schonfeld: A new parallel message-distribution technique for cost-based steganography, 2017.

Patents

3. S. Deshmukh, A. Dasgupta, **C. Agarwal**, B. Krishnamurthy, G. Theocharous, J. Subramanian.: Novel Trajectory-based Explainability Framework for RL-based Decision Making. Internal Reference: P11853-US.

2. M. Hemani, A. Seth, **C. Agarwal**: Debiasing vision-language models with additive residual learning. Internal Reference: P11919-US.

1. T. Menta, A. Patil, S. Jandial, Balaji K, **C. Agarwal**, M. Sarkar: HASTE: A Novel Method and Apparatus to Estimate Transferability using Hard Subsets. Internal Reference: P11683-US.

Teaching Experience

Guest Lecture at Harvard University   Spring 2021, 2023
Course on Interpretability and Explainability in Machine Learning

Teaching Assistant
University of Illinois at Chicago   Spring, Fall 2014 - 2020
Pattern Recognition, Image Analysis & Computer Vision,
Digital Signal Processing, Neural Networks.

Tutorials

Explainability in Graph Deep Learning for Biomedicine   ISMB 2024
Training the Next-Generation of AI Students   Excellence School 2023
Explainable ML in the Wild: When Not to Trust Your Explanations   FAccT 2021

Workshop

Workshop on Regulatable ML   NeurIPS 2023

Invited Talks

Computer Vision Talks   2023
TrustML Young Scientists Seminars at RIKEN-AIP, Japan   2022
Adobe Research: XAI: Challenges and Solutions   2022
CAI Summer School at IIIT-Delhi   2022
LOGML Summer School   2022
2d3d.ai   2021
W&B - Weights & Biases Salon   2020

Mentorship

Current Advisee
Nicholas Kroeger, Ph.D. Student, University of Florida   2023-Present
Dan Ley, Ph.D. Student, Harvard University   2023-Present
Satyapriya Krishna, Ph.D. Student, Harvard University   2020-Present
Elita Lobo, Ph.D. Student, University of Massachusetts, Amherst   2023-Present
Sree Harshan Tanneru, Masters Student, Harvard University   2023-Present
Nikhil Nayak, Masters Student, Harvard University   2023-Present
Surgan Jandial, Research Engineer, Adobe   2023-Present
Abhinav Java, Research Engineer, Adobe   2023-Present

Past Advisee and Interns
Martin Pawelczyk, Ph.D. Student, University of Tübingen   2021-2022
Valentina Giunchiglia, Ph.D. Student, Imperial College London   2022-2023
Chirag Varun Shukla, Ph.D. Student, LMU Munich   2022-2023
Jiali Cheng, Ph.D. Student, University of Massachusetts Lowell   2022-2023

Ashish Seth, Masters Student, IIT Madras   2022-2023
Owen Queen, Undergrad, University of Tennessee, Knoxville   2021-2022
Tarun R Menta, Undergrad, IIT Hyderabad   2022-2023
Nari Johnson, Undergrad, Harvard University   2022
Eshika Saxena, Undergrad, Harvard University   2022
Isha Puri, Undergrad, Harvard University   2022
Daniel D'souza, Data Scientist, Proquest   2021-2022
Shripad V Deshmukh, Research Engineer, Adobe   2022-2023

Community Service

Founder: Agyeya Artificial IQ Foundation	2023-Present
Open Collaboration Initiatives: TrustworthyML Initiative and MLCollective	2021-2023
External Ph.D. Examiner: Jessica Rumbelow - University of St. Andrews	2023

Program Committee for Workshops:

RegML - Regulatable Machine Learning (RegML)	NeurIPS, 2023
XAI4CV - Explainable AI for Computer Vision (XAI4CV) Workshop	CVPR, 2023
SRML - Workshop on Socially Responsible Machine Learning	ICLR, 2022
AdvML - New Frontiers in Adversarial Machine Learning	ICML, 2022
SRML - Workshop on Socially Responsible Machine Learning	ICML,2021
SeSML - Workshop on Security and Safety in Machine Learning Systems	ICLR, 2021
AROW - Workshop on Adversarial Robustness in the Real World	ECCV, 2020-2021
WHI - Workshop on Human Interpretability in Machine Learning	ICML, 2020

Program Committee for Conferences:

NeurIPS - Advances in Neural Information Processing Systems	2021-2023
NeurIPS - Datasets and Benchmark Track	2022-2023
KDD - ACM SIGKDD Conference on Knowledge Discovery and Data Mining	2021-2023
ICML - International Conference on Machine Learning	2021-2023
FAccT - ACM Conference on Fairness, Accountability, and Transparency	2022-2023
ICLR - International Conference on Learning Representations	2022-2023
AAAI - AAAI International Conference on Artificial Intelligence	2023-2024
AISTATS - International Conference on Artificial Intelligence and Statistics	2023
WACV - IEEE/CVF Winter Conference on Applications of Computer Vision	2023
CVPR - IEEE/CVF Conference on Computer Vision and Pattern Recognition	2023
ICCV - IEEE/CVF International Conference on Computer Vision	2023
ACL - ACL Rolling Review	2023
LOG - Learning on Graphs Conference	2022

Journal Reviewing:

TMLR - The Transactions on Machine Learning Research	2022-2023
TMI - IEEE Transactions on Medical Imaging	2022