

Chirag Agarwal

[Google Scholar](#)

Email: chiragagarwal@virginia.edu

Phone: (865)-406-2887

ACADEMIC & PROFESSIONAL EXPERIENCE

University of Virginia

Assistant Professor

2024 – Present

School of Data Science (primary appointment)

Department of Computer Science (by courtesy)

School of Medicine, Radiology and Medical Imaging (by courtesy)

[Aikyam Lab](#)

Harvard University

Postdoctoral Research Fellow

2020–2022, 2023–2024

Host: [Prof. Hima Lakkaraju](#) and [Prof. Marinka Zitnik](#)

Adobe

Research Scientist

2022 – 2023

Auburn University

Research Assistant

Summer 2019

Robert Bosch LLC

Computer Vision/Augmented Reality Intern

Summer 2018

Tempus labs Inc.

Imaging Science Intern

Spring 2018

Kitware Inc.

Research and Development Intern

Summer 2017

Geisinger Health Systems

Research Intern

Summer 2016

EDUCATION

University of Illinois at Chicago

Ph.D. in Electrical and Computer Engineering

2020

– Thesis: *Robustness and Explainability of Deep Neural Networks*

– Committee: [Prof. Dan Schonfeld](#), [Prof. Bharati Prasad](#), [Prof. Mojtaba Soltanalian](#),
[Prof. Piotr Gmytrasiewicz](#), [Prof. Anh Nguyen](#)

University of Illinois at Chicago

M.S. in Electrical and Computer Engineering

2018

SELECTED HONORS & ACHIEVEMENTS

[Baugh Family Fellow](#)

2026

[AAAI New Faculty Highlight](#)

2026

SDS Pilot Grant

2025

[Capital One Fellowship](#)

2025

Tinker Research Grant

2025

OpenAI Researcher Access Program Award

2025

The Environmental Institute’s Colab Award	2025
Cohere For AI Research Grant	2025
The LaCross AI Institute Fellowship in AI Research	2025
Dean’s Strategic Fund	2024
Top Reviewer for NeurIPS	2023
Adobe Data Science Research Award	2023
Spotlight presentation , NeurIPS Ro-FoMo Workshop in Foundation Models	2023
Spotlight paper , ICML	2021
AINet Fellow by DAAD	2021
Spotlight presentation , ICML workshop on Human Interpretability in Machine Learning	2020
Spotlight paper , IEEE Conference on Image Processing (ICIP)	2019

SELECTED GRANTS & AWARDS

Baugh Family Fellow(US \$5,000) – PI	2026
SDS Pilot Grant (US \$20,0000) – PI	2025
Capital One Fellowship (US \$128,500) – Sole PI	2025
Tinker Research Grant (US \$5,000) – Sole PI	2025
OpenAI Researcher Access Program Award (US \$1,000) – Sole PI	2025
The Environmental Institute’s Colab Award (US \$100,000) – PI	2025
Cohere For AI Research Grant (US \$2,000) – Sole PI	2025
The LaCross AI Institute Fellowship in AI Research (US \$100,000) – PI	2025
Dean’s Strategic Fund (US \$7,440) – PI	2024
Adobe Data Science Research Award (US \$50,000) – co-PI	2023
Harvard Data Science Initiative Microsoft Azure Credits (US \$22,224) – co-PI	2023
AI for Social Good Google Workshop (US \$10,000) – co-PI	2021
2 × Research Proposal accepted by Google Cloud Platform (US \$1,000) – Sole PI	2020

RESEARCH ARTICLES

† denotes the author I co-mentored with the PI; * indicates an equal contribution.

Articles in Peer-Reviewed Journals

78. **C. Agarwal**, O. Queen†, H. Lakkaraju, M. Zitnik: Evaluating Explainability for Graph Neural Networks, *Nature Scientific Data*, 2023.
77. H. Honarvar, **C. Agarwal**, S. Somani, A. Vaid, J. Lampert, T. Wanyan, V. Y. Reddy, G. N. Nadkarni, R. Miotto1, M. Zitnik, F. Wang, B. S. Glicksberg: Enhancing convolutional neural network predictions of electrocardiograms with left ventricular dysfunction using a novel sub-waveform representation, *Cardiovascular Digital Health Journal*, 2022.
76. **C. Agarwal**, S. Gupta†, M. Y. Najjar, T. E. Weaver, X. J. Zhou, D. Schonfeld, B. Prasad: Deep Learning Analyses of Brain MRI to Identify Sleepiness in Treated Obstructive Sleep Apnea: A Pilot Study, *Journal of Sleep and Vigilance (JSV)*, 2022.
75. B. Prasad*, **C. Agarwal***, E. Schonfeld, D. Schonfeld, B. Mokhlesi: Deep learning applied to polysomnography to predict blood pressure in obstructive sleep apnea and obesity hypoventilation: A proof-of-concept study, *Journal of Clinical Sleep Medicine (JCSM)*, 2020.
74. **C. Agarwal**, J. Klobusicky, D. Schonfeld: Convergence of backpropagation with momentum for network architectures with skip connections, *Journal of Computational Mathematics (JCM)*, 2019.

73. E. Cha, Y. Veturi, **C. Agarwal**, M. Arbabshirani, S. Pendergrass: Using Adipose Measures from Electronic Health Record Imaging Based Data for Discovery, *Journal of Obesity*, 2018.

Articles in Peer-Reviewed Conference Proceedings

72. A. Ghosh, S. Sridhar, R.K. Ravi, M. Muhsin, S. Saha, **C. Agarwal**: CLINIC: Evaluating Multilingual Trustworthiness in Language Models for Healthcare, *ICML*, 2026.
71. T. B Higgins, A. Mamalakis, **C. Agarwal**: SynopticBench: Evaluating Vision-Language Models on Generating Weather Forecast Discussions of the Future, *Climate Informatics*, 2026. **(Oral Presentation)**
70. E. Onyame, A. Ghosh, S. Baidya, S. Saha, X. Chen, **C. Agarwal**: CURE-Med: Curriculum-Informed Reinforcement Learning for Multilingual Medical Reasoning, *ACL*, 2026. **(Oral Presentation)**
69. A. Saiyed, S. Sadiekh, **C. Agarwal**: Towards Understanding the Robustness of Sparse Autoencoders, *ACL*, 2026.
68. S. Petkar, A. Vempati, A. Sinha, P. Kumarauguru, **C. Agarwal**: A Graph Talks, But Who's Listening? Rethinking Evaluations for Graph-Language Models, *ACL*, 2026. **(New Perspectives in Graph Machine Learning Workshop, NeurIPS'25)**
67. J. Cheng, Z. Chen, **C. Agarwal**, H. Amiri: Toward Understanding Unlearning Difficulty: A Mechanistic Perspective and Circuit-Guided Difficulty Metric, *ACL*, 2026.
66. **C. Agarwal**: Towards Trustworthy Multimodal AI Systems, *AAAI*, 2026. **(Oral Presentation)**
65. S. Sabrina, E. Ericheva, **C. Agarwal**: Polarity-Aware Probing for Quantifying Latent Alignment in Language Models, *AAAI*, 2026. **(Oral Presentation)**
64. R. Zhou, G. Nguyen, N. Kharya, A. Nguyen, **C. Agarwal**: Improving Human Verification of LLM Reasoning through Interactive Explanation Interfaces, *IUI*, 2026.
63. A. Ghosh, D. Datta, S. Saha, **C. Agarwal**: The Multilingual Mind: A Survey of Multilingual Reasoning in Language Models, *EMNLP*, 2025.
62. A. Seth, U. Tyagi, R. Selvakumar, N. Anand, S. Kumar, S. Ghosh, R. Duraiswami, **C. Agarwal**, D. Manocha: EGOILLUSION: Benchmarking Hallucinations in Egocentric Video Understanding, *EMNLP*, 2025.
61. T. Menta, S. Agrawal, **C. Agarwal**: Analyzing Memorization in Large Language Models through the Lens of Model Attribution, *NAACL*, 2025. **(Oral Presentation)**
60. A. Java, S. Shahid, **C. Agarwal**: Towards Operationalizing Right to Data Protection, *NAACL*, 2025.
59. E Lobo†, **C. Agarwal**, H Lakkaraju: On the Impact of Fine-Tuning on Chain-of-Thought Reasoning, *NAACL*, 2025.
58. T. Han, A. Kumar, **C. Agarwal**, H. Lakkaraju: Towards Safe and Aligned Large Language Models for Medicine, *NeurIPS*, 2024. **(ICML Next Generation of AI Safety Workshop, 2024)**
57. A. Kumar, **C. Agarwal**, S. Srinivas, S. Feizi, H. Lakkaraju: Certifying LLM Safety against Adversarial Prompting, *COLM*, 2024. **(Featured in Science News Magazine and D³ Institute at Harvard)**
56. S. Krishna†, **C. Agarwal**, H. Lakkaraju: On the Impact of Adversarially Robust Models on Algorithmic Recourse, *AIES*, 2024.
55. S. Krishna†, **C. Agarwal**, H. Lakkaraju: Understanding the Effects of Iterative Prompting on Truthfulness, *International Conference on Machine Learning (ICML)*, 2024.
54. S. H. Tanneru†, **C. Agarwal**, H. Lakkaraju: Uncertainty In Explanations Of Large Language Models, *International Conference on Artificial Intelligence and Statistics (AISTATS)*, 2024. **Spotlight Presentation at the NeurIPS R0-FoMo Workshop, 2023**
53. M. Llorde, D. Ganguly, S. Bhatia, **C. Agarwal**: Explain like I am BM25: Interpreting a Dense Model's Ranked-List with a Sparse Approximation, *ACM SIGIR Conference on Research and Development in Information Retrieval (SIGIR)*, 2023.
52. A. Seth, M. Hemani, **C. Agarwal**: DeAR: Debiasing Vision-Language Models with Additive Residuals, *Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023.

51. S. Deshmukh†, A. Dasgupta, B. Krishnamurthy, N. Jiang, **C. Agarwal**, J. Subramanian, G. Theodorou: Trajectory-based Explainability Framework for Offline RL, *International Conference on Learning Representations (ICLR)*, 2023.
50. J. Cheng†, G. Dasoulas, H. He, **C. Agarwal**, M. Zitnik: GNNDelete: A General Unlearning Strategy for Graph Neural Networks, *International Conference on Learning Representations (ICLR)*, 2023.
49. V. Giunchiglia, C. V. Shukla, G. Gonzalez, **C. Agarwal**: Towards Training GNNs using Explanation Directed Message Passing, *Proceedings of the First Learning on Graphs Conference (LoG)*, 2022.
48. **C. Agarwal**, E. Saxena†, S. Krishna†, M. Pawelczyk†, N. Johnson†, I. Puri†, M. Zitnik, H. Lakkaraju: OpenXAI: Towards a Transparent Evaluation of Model Explanations, *Conference on Neural Information Processing Systems (NeurIPS)*, 2022.
47. **C. Agarwal**, D. D’Souza†, S. Hooker: Estimating Example Difficulty using Variance of Gradients, *Conference on Computer Vision and Pattern Recognition (CVPR)*, 2022.
46. **C. Agarwal**, M. Zitnik, H. Lakkaraju: Probing GNN Explainers: A Rigorous Theoretical and Empirical Analysis of GNN Explanation Methods, *International Conference on Artificial Intelligence and Statistics (AISTATS)*, 2022.
45. M. Pawelczyk†, **C. Agarwal**, S. Joshi, S. Upadhyay, H. Lakkaraju: Exploring Counterfactual Explanations Through the Lens of Adversarial Examples: A Theoretical and Empirical Analysis, *International Conference on Artificial Intelligence and Statistics (AISTATS)*, 2022.
44. **C. Agarwal**, H. Lakkaraju, M. Zitnik: Towards a Unified Framework for Fair and Stable Graph Representation Learning, *Conference on Uncertainty in Artificial Intelligence (UAI)*, 2021.
43. S. Agarwal, S. Jabbari, **C. Agarwal**, S. Upadhyay, Z. S. Wu, H. Lakkaraju: Towards the Unification and Robustness of Perturbation and Gradient Based Explanations, *International Conference on Machine Learning (ICML)*, 2021. **(Spotlight Presentation)**
42. **C. Agarwal***, S. Khobahi*, D. Schonfeld, M. Soltanian: CoroNet: A Deep Network Architecture for Semi-Supervised Task-Based Identification of COVID-19 from Chest X-ray Images, *SPIE Medical Imaging*, 2021.
41. **C. Agarwal**, A. Nguyen: Explaining image classifiers by removing input features using generative models, *Asian Conference on Computer Vision (ACCV)*, 2020.
40. N. Bansal*, **C. Agarwal***, A. Nguyen*: SAM: The Sensitivity of Interpretability Methods to Hyperparameters, *Conference on Computer Vision and Pattern Recognition (CVPR)*, 2020. **(Oral presentation)**
39. **C. Agarwal**, S. Khobahi, A. Bose, M. Soltanian, D. Schonfeld: Deep-URL: A Model-Aware Approach To Blind Deconvolution Based On Deep Unfolded Richardson-Lucy Network, *IEEE Conference on Image Processing (ICIP)*, 2020.
38. **C. Agarwal**, A. Nguyen, D. Schonfeld: Improving Robustness to Adversarial Examples by Encouraging Discriminative Features, *IEEE Conference on Image Processing (ICIP)*, 2019. **(Spotlight presentation)**
37. M. Aloraini, M. Sharifzadeh, **C. Agarwal**, D. Schonfeld: Statistical Sequential Analysis for Object-based Video Forgery Detection, *Electronic Imaging*, 2019.
36. N. Khobragade*, **C. Agarwal***: Multi-class segmentation of neuronal electron microscopy images using deep learning, *SPIE Medical Imaging*, 2018.
35. **C. Agarwal**, M. Sharifzadeh, D. Schonfeld: CrossEncoders: A complex neural network compression framework, *IS&T International Symposium on Electronic Imaging*, 2018.
34. M. Sharifzadeh, **C. Agarwal**, M. Aloraini, D. Schonfeld: Convolutional neural network steganalysis’s application to steganography, *IEEE Visual Communications and Image Processing (VCIP)*, 2017.
33. **C. Agarwal**, A.H. Dallal, M.R. Arbabshirani, A. Patel, G. Moore: Unsupervised quantification of abdominal fat from CT images using Greedy Snakes, *SPIE Medical Imaging*, 2017.
32. A.H. Dallal, **C. Agarwal**, M.R. Arbabshirani, A. Patel, G. Moore: Automatic estimation of heart boundaries and cardiothoracic ratio from chest X-ray images, *SPIE Medical Imaging*, 2017.
31. M.R. Arbabshirani, A.H. Dallal, **C. Agarwal**, A. Patel, G. Moore: Accurate segmentation of lung fields on chest radiographs using deep convolutional networks, *SPIE Medical Imaging*, 2017.

30. **C. Agarwal**, A. Bose, S. Maiti, N. Islam, S.K. Sarkar: Enhanced data hiding method using DWT based on Saliency model, *IEEE International Conference on Signal Processing, Computing and Control (ISPPC)*, 2013.
29. S. Maiti, **C. Agarwal**, A. Bose, S.K. Sarkar: Robust data hiding technique in wavelet domain using saliency map, *International Journal of Advances in Engineering and Technology*, 2013.
28. N. Islam S. Maiti, A. Bose, **C. Agarwal**, S. K. Sarkar: An Improved Method of Pre-Filter Based Image Watermarking in DWT Domain, *International Journal of Computer Science and Technology*, 2013.

Preprints and Workshop Articles

27. S. Ghosh, R. Venkatesh Babu, **C. Agarwal**: Do vision language models need to process image tokens?, *TRUE-V @ CVPR*, 2026. **(Oral Presentation)**
26. S. Belsare, S. Nikhil, S. Kumar, P. Kumaraguru, **C. Agarwal**: GridVQA-X: A Framework for Evaluating Multimodal Explainability Methods, *TRUE-V @ CVPR*, 2026.
25. E. Lobo, X. Chen, J. Meng, N. Xi, Y. Jiao, **C. Agarwal**, Y. Zick, and Y. Gao. STRUCTUREDAGENT: Planning with AND/OR Trees for Long-Horizon Web Tasks, *arXiv*, 2026. **(Amazon Collaboration)**
24. D. Zhang, S. Betala, **C. Agarwal**. Quantifying Explanation Quality in Graph Neural Networks using Out-of-Distribution Generalization, *arXiv*, 2026.
23. T. Nguyen, T. Truong, R. Zhou, T. Bui, **C. Agarwal**, A. T. Nguyen: PageGuide: Browser extension to assist users in navigating a webpage and locating information, *arxiv*, 2026.
22. J. Cheng, **C. Agarwal**, H. Amiri: Do Students Debias Like Teachers? On the Distillability of Bias Mitigation Methods, *arxiv*, 2025.
21. S. Wang, Z. Lei, Z. Tan, J. Li, J. Rasero, A. Zhang, **C. Agarwal**: Interpretable Neuropsychiatric Diagnosis via Concept-Guided Graph Neural Networks, *arXiv*, 2025.
20. N. Safari, I. Levy, A. Zhang, **C. Agarwal**, J. M Cortes, K. A Pelphrey, J. Darrell Van Horn, J. Rasero: Brain morphology network alterations in adolescents with autism spectrum disorder: a sex-stratified study, *arXiv*, 2025.
19. **C. Agarwal**: Rethinking Explainability in the Era of Multimodal AI, *arXiv*, 2025
18. A Seth, D Manocha, **C. Agarwal**: Towards a Systematic Evaluation of Hallucinations in Large-Vision Language Models, *EMNLP*, 2025.
17. D. Ley†, S. H. Tanneru†, **C. Agarwal**, H. Lakkaraju: On the Difficulty of Faithful Chain-of-Thought Reasoning in Large Language Models, *ICML TiFA Workshop*, 2024. **(Featured in OpenAI o1 System Card Report)**
16. **C. Agarwal**, S. H. Tanneru, H. Lakkaraju: Faithfulness vs. Plausibility: On the (Un)Reliability of Explanations from Large Language Models, *arXiv*, 2024. **(Featured in OpenAI o1 System Card Report)**
15. N. Kroeger†, D. Ley†, S. Krishna†, **C. Agarwal**, H. Lakkaraju: Are Large Language Models Post Hoc Explainers?, *Preliminary version presented at the NeurIPS XAIA Workshop*, 2023.
14. A. Java, S. Jandial, **C. Agarwal**: Towards Fair Knowledge Distillation using Student Feedback, *Preliminary version presented at the Efficient Systems for Foundation Models, ICML 2023*.
13. S.V. Deshmukh, Srivatsan R, S. Vijay, J. Subramanian, **C. Agarwal**: Counterfactual Explanation Policies in RL, *Preliminary version presented at "Could it have been different?" Counterfactuals in Minds and Machines Workshop, ICML 2023*.
12. T. R. Menta†, S. Jandial†, A. Patil, Vimal KB, S. Bachu, B. Krishnamurthy, V. N. Balasubramanian, **C. Agarwal**, M. Sarkar: Towards Estimating Transferability using Hard Subsets, *arXiv*, 2023.
11. **C. Agarwal**: Intriguing Properties of Visual-Language Model Explanations, *Preliminary version presented at RTML Workshop, ICLR 2023*.
10. S. Krishna†, **C. Agarwal**, H. Lakkaraju: On the Impact of Adversarially Robust Models on Algorithmic Recourse, *Preliminary version presented at Trustworthy and Socially Responsible ML Workshop, NeurIPS 2022*.
9. **C. Agarwal**, N. Johnson†, M. Pawelczyk†, S. Krishna†, E. Saxena†, M. Zitnik, H. Lakkaraju: Rethinking Stability for Attribution-based Explanations, *Preliminary version presented at PAIR² Struct Workshop, ICLR, 2022. (Oral Presentation)*

8. D. D'Souza†, Z. Nussbaum†, **C. Agarwal**, S. Hooker: A Tale Of Two Long Tails, *Preliminary version presented at Uncertainty & Robustness in Deep Learning Workshop, ICML, 2021*.
7. H. Honarvar, **C. Agarwal**, S. Somani, A. Vaid, J. Lampert, T. Wanyan, V. Y. Reddy, G. N. Nadkarni, R. Miotto1, M. Zitnik, F. Wang, B. S. Glicksberg: A novel representation of electrocardiogram waveforms for enhancing deep learning predictions, *Preliminary version presented at Interpretable Machine Learning in Healthcare Workshop, ICML, 2021*.
6. **C. Agarwal***, P. Chen*, A. Nguyen: Intriguing generalization and simplicity of adversarially trained neural networks, *Preliminary version presented at Human Interpretability in Machine Learning Workshop, ICML, 2020. (Spotlight Presentation)*
5. **C. Agarwal**, B. Dong, D. Schonfeld, A. Hoogs: An explainable adversarial robustness metric for deep learning neural networks, 2018.
4. M. Sharifzadeh, **C. Agarwal**, M. Salarian, D. Schonfeld: A new parallel message-distribution technique for cost-based steganography, 2017.

Patents

3. T. Menta, A. Patil, S. Jandial, Balaji K, **C. Agarwal**, M. Sarkar: Systems and methods for machine learning transferability. Application number: 18178225, 2024
2. M. Hemani, A. Seth, **C. Agarwal**: Debiasing vision-language models with additive residuals. Application number: 18322253, 2024
1. S. Deshmukh, A. Dasgupta, **C. Agarwal**, B. Krishnamurthy, G. Theocharous, J. Subramanian.: Novel Trajectory-based Explainability Framework for RL-based Decision Making. Internal Reference: P11853-US.

TEACHING EXPERIENCE

Instructor, Explainable AI <i>School of Data Science, UVA</i>	Spring 2026
Instructor, Decoding Large Language Models <i>School of Data Science, UVA</i>	Spring 2025
Guest Lecture on Explainable Artificial Intelligence <i>Course on Foundation of Data Science, UVA</i>	Fall 2024-2025, Spring 2025
Guest Lecture on Explainable Artificial Intelligence <i>Course on Interpretability and Explainability in ML, Harvard University</i>	Spring 2021, 2023
Teaching Assistant University of Illinois at Chicago <i>Pattern Recognition, Image Analysis & Computer Vision, Digital Signal Processing, Neural Networks.</i>	Spring, Fall 2014 - 2020

TUTORIALS

Explainability in Graph Deep Learning for Biomedicine	ISMB 2024
Training the Next-Generation of AI Students	Excellence School 2023
Explainable ML in the Wild: When Not to Trust Your Explanations	FACCT 2021

WORKSHOP

Generative AI meets Explainable AI	XAI 2025
Workshop on Regulatable ML	NeurIPS 2023-2025

INVITED TALKS

World Health Organization (WHO) South-East Asia Office Webinar Series	2026
Charles L. and Ann Brown Distinguished Colloquium series @ UVA	2026
AIxHealth seminar at GeorgiaTech	2026
Towards Trustworthy Multimodal AI Systems @ AAAI	2026
Alignment workshop at FAR.AI	2025
Multimodal Reasoning Workshop at ACM SIGMM	2025
The AI and Machine Learning Seminar @ UVa	2025
UMass Lowell Computer Science Colloquia	2025
International Institute of Information Technology - Hyderabad (IIIT-H)	2025
Social Contagions, Artificial Intelligence, and Democracy Workshop	2025
FAR.AI	2024
UVA Conference on Leadership in Business, Data and Intelligence – LaCross AI Institute	2024
Workshop on Privacy and Interpretability in Generative AI: Peering into the Black Box	2024
Computer Vision Talks	2023
TrustML Young Scientists Seminars at RIKEN-AIP, Japan	2022
Adobe Research: XAI: Challenges and Solutions	2022
CAI Summer School at IIIT-Delhi	2022
LOGML Summer School	2022
2d3d.ai	2021
W&B - Weights & Biases Salon	2020

MENTORSHIP

Current Advisee

Timothy Higgins, Postdoctoral Research Associate, UVA	2025-Present
Eleanor LaRocco, Ph.D. Student, UVA	2025-Present
Eric Onyame, Ph.D. Student, UVA	2025-Present
Ahson Saiyed, Ph.D. Student, UVA	2025-Present
Madelyn Mathai, Ph.D. Student, UVA	2025-Present
Runtao Zhou, Ph.D. Student, UVA	2025-Present
Ding Zhang, Ph.D. Student, UVA	2025-Present
Yuyang Cheng, Ph.D. Student, UVA	2025-Present
Akash Ghosh, Ph.D. Student, IIT Patna	2024-Present
Nishitha Khasnavis, Undergraduate student, UVA	2025-Present
Dailin Li, Undergraduate student, UVA	2025-Present

Past Advisee and Interns

Elita Lobo, Ph.D. Student, UMass Amherst → Amazon	2023-2026
Ashish Seth, Ph.D. Student, UMD	2024-2025
Susmit Agrawal, MS Student, IIT-H → PhD Student @ Tübingen (IMPRS-IS)	2024-2025
Tarun R Menta, Research Engineer @ Adobe → DataLab	2024-2025
Abhinav Java, Research Engineer @ Adobe → Microsoft	2023-2025
Simra Sahid, Research Engineer @ Adobe → Microsoft	2024-2025
Dan Ley, Ph.D. Student, Harvard University	2023-2024
Nicholas Kroeger, Ph.D. Student, University of Florida	2023-2024
Sree Harsha Tanneru, MS Student @ Harvard → Research Engineer, DeepMind	2023-2024
Satyapriya Krishna, Ph.D. Student, Harvard University → Amazon	2020-2024
Martin Pawelczyk, Ph.D. Student, University of Tübingen	2021-2022
Valentina Giunchiglia, Ph.D. Student, Imperial College London	2022-2023
Chirag Varun Shukla, Ph.D. Student, LMU Munich	2022-2023
Jiali Cheng, Ph.D. Student, University of Massachusetts Lowell	2022-2023
Surgan Jandial, Research Engineer @ Adobe → MS Student @ CMU	2022-2023

Shripad V Deshmukh, Research Engineer @ Adobe → Ph.D. student at UMass Amherst	2022-2023
Nari Johnson, Undergrad @ Harvard University → Ph.D. student @ CMU	2022
Eshika Saxena, Undergrad, Harvard University → Meta	2022
Isha Puri, Undergrad, Harvard University → Ph.D. student @ MIT	2022
Owen Queen, Undergrad, University of Tennessee, Knoxville → Ph.D. student @ Stanford	2021-2022
Daniel D'souza, Data Scientist, Proquest → Cohere	2021-2022

COMMUNITY SERVICE

Founder: Agyeya Artificial IQ Foundation	2023-Present
Open Collaboration Initiatives: TrustworthyML Initiative and MLCollective	2021-2023
External Ph.D. Examiner:	
Anna Hedström - TU Berlin	2025
Sichao Li - Australian National University	2025
Jessica Rumbelow - University of St. Andrews	2023
Reviewer for Grants and Proposals:	
Deutsche Forschungsgemeinschaft (DFG) German Research Foundation	2025
Precision Health Initiative Collaborative Faculty Grant UVA	2025
Area Chair/Senior Program Committee:	
AAAI Alignment Track - <i>AAAI International Conference on Artificial Intelligence</i>	2026
ICLR - <i>International Conference on Learning Representations</i>	2026
NeurIPS - <i>Advances in Neural Information Processing Systems</i>	2026
Program Committee for Conferences:	
ICML - <i>International Conference on Machine Learning</i>	2021-2026
COLM - <i>Conference on Language Modeling</i>	2026
NeurIPS - <i>Advances in Neural Information Processing Systems</i>	2021-2025
NeurIPS - <i>Datasets and Benchmark Track</i>	2022-2025
KDD - <i>ACM SIGKDD Conference on Knowledge Discovery and Data Mining</i>	2021-2023
FAccT - <i>ACM Conference on Fairness, Accountability, and Transparency</i>	2022-2023
ICLR - <i>International Conference on Learning Representations</i>	2022-2025
AAAI - <i>AAAI International Conference on Artificial Intelligence</i>	2023-2026
AIES - <i>AAAI Conference on AI, Ethics, and Society</i>	2024-2025
XAI World Conference	2024
AISTATS - <i>International Conference on Artificial Intelligence and Statistics</i>	2023
WACV - <i>IEEE/CVF Winter Conference on Applications of Computer Vision</i>	2023
CVPR - <i>IEEE/CVF Conference on Computer Vision and Pattern Recognition</i>	2023,2026
ICCV - <i>IEEE/CVF International Conference on Computer Vision</i>	2023
ACL - <i>ACL Rolling Review</i>	2023
LOG - <i>Learning on Graphs Conference</i>	2022
Program Committee for Workshops:	
ACL-SRW - <i>Student Research Workshop (SRW)</i>	ACL, 2025
RegML - <i>Regulatable Machine Learning (RegML)</i>	NeurIPS, 2023-2024
WHI - <i>Workshop for Women in Machine Learning (WiML)</i> ,	NeurIPS, 2024
XAI4CV - <i>Explainable AI for Computer Vision (XAI4CV) Workshop</i>	CVPR, 2023
SRML - <i>Workshop on Socially Responsible Machine Learning</i>	ICLR, 2022
AdvML - <i>New Frontiers in Adversarial Machine Learning</i>	2022, 2024
SRML - <i>Workshop on Socially Responsible Machine Learning</i>	ICML,2021

[SeSML](#) - *Workshop on Security and Safety in Machine Learning Systems*

ICLR, 2021

[AROW](#) - *Workshop on Adversarial Robustness in the Real World*

ECCV, 2020-2021

[WHI](#) - *Workshop on Human Interpretability in Machine Learning*

ICML, 2020

Journal Reviewing:

Transactions on Pattern Analysis and Machine Intelligence

2025

Nature Communications

2024

TMLR - The Transactions on Machine Learning Research

2022-2024

TMI - IEEE Transactions on Medical Imaging

2022